



Best Practices for PL/SQL

Objectives



At the end of this session, you will be able to:

- Know the Best Practices for PL/SQL Program
- Understand Modularity
- Learn Challenges of Local Module
- Learn Anchor Declarations of Variables
- Know the Use of Bind Variables
- Learn the Use of Cursors & Control Structures
- Use Bulk Collect Clause
- Use the FORALL Clause
- Know BEFORE VS AFTER TRIGGER
- Understand Optimizing SQL

- Best Practices for PL/SQL Program
- Modularity
- Challenges of Local Module
- Learn Anchor Declarations of Variables
- Use of Bind Variables
- Cursors & Control Structures
- Bulk Collect Clause
- The FORALL Clause
- BEFORE VS AFTER TRIGGER
- Optimizing SQL

- A best practice workflow for single program construction
 - Focus on testing and improving the readability of one's code

- A best practice workflow for the application as a whole
 - Focus on how best to write (and not write) SQL

- Improves software Reusability
- Hides the Complexity of execution of a specific operation behind a name
- We use top-down design to hide Complexity
- We use small and narrowly-focused packages
 - Encapsulate logically related types, items & subprograms in a named PL/SQL module
- Advantages of using Packages:
 - Easier Application Design
 - Overloading
 - Information Hiding
 - Better Performance

Always:

- Reuse existing programs whenever possible
- Write tiny chunks (blocks) of code
- Build lots of Local or Nested Modules
 - Avoid Spaghetti code
 - Keep your executable sections small / tiny

- Requires Discipline
 - Always be on the lookout for opportunities to Refactor
 - Make conscious decisions about the scope of your programs
- Need to Read from the Bottom, Up
 - For a program with local modules, start at the Bottom (main executable section) and read Up
 - Takes some time getting used to
- Your IDE should reveal the Internal Program Structure
 - Toad's Program Navigator & SQL Navigator's Code Explorer are ideal facilitators of Top-down Design and Local Modularization

- Define Subprograms close to usage
- A subprogram can be defined at any of these levels:
 - Local - Within another Subprogram
 - Private - Defined in the Package Body
 - Public - Defined in the Package Specification
- The best rule to follow is:

Define subprograms as close as possible to their usage.

The shorter the distance from usage to definition, the easier it is to find, understand and maintain that code

- 2 choices to declare a variable:
 - Hard-coding the data type
 - Anchoring the data type to another Structure

Hard-coded Declaration

```
ename VARCHAR2(30);  
totsales NUMBER (10,2);
```

Anchored Declarations

```
v_ename emp.ename%TYPE;  
Totsales pkg.sales_amt%TYPE;
```

- Prefer Anchored Declarations to Explicit data type references
 - %TYPE for Scalar Structures
 - %ROWTYPE for Composite Structures

- Avoids unnecessary Parsing
- A key to improved Code Performance
- PL/SQL code improves by several orders of magnitude, depending on the Type and Size of data

```
DECLARE
BEGIN
  FOR i IN 1 .. 1000 LOOP
    EXECUTE IMMEDIATE
      'SELECT item_id, qty
      FROM items
      WHERE qty = ' || i;
  END LOOP;
END;
```

```
DECLARE
BEGIN
  FOR i IN 1 .. 1000 LOOP
    EXECUTE IMMEDIATE
      'SELECT item_id, qty
      FROM items
      WHERE qty = :x' using i;
  END LOOP;
END;
```

- Use of Explicit cursors for fetching data from a table
 - Faster
 - Exception Handling is easier
 - Greater programmatic control

- Use of Cursor FOR LOOP
 - No need to write OPEN, FETCH and CLOSE
 - Results in lesser Lines of Code & Errors

- An Exception is not raised if no rows are affected by an UPDATE or DELETE statement

SQL%ROWCOUNT or SQL%FOUND is the only way to trap such kind of conditions

- Select Your Control Structures carefully

Multiple Conditional Statements

```
IF quantity <= 5 THEN  
  Shipping_method :=standard_shipping;  
END IF;
```

```
IF quantity > 5 THEN  
  Shipping_method :=  
expedited_shipping;  
END IF;
```

Replacing IF with ELSIF

```
IF quantity <= 5 THEN  
  Shipping_method :=  
standard_shipping;  
ELSIF quantity > 5 THEN  
  Shipping_method :=  
expedited_shipping;  
END IF;
```

- A better alternative to Nested IF-ELSE statements
- Handy when number of conditions to check for is large
- Using CASE for mutually exclusive condition checking:

```
CASE quantity  
    WHEN <=5 THEN shipping_method := standard_shipping;  
    WHEN > 5 THEN shipping_method := expedited_shipping;  
    ELSE shipping_method := 0  
END CASE;
```

Fetch into Cursor Records

```
name VARCHAR2 (30);  
minbal NUMBER(10,2);  
BEGIN  
OPEN company_pkg.allrows;  
FETCH company_pkg.allrows  
INTO name, minbal;  
IF name = 'ACME' THEN ...  
CLOSE company_pkg.allrows;
```

Fetching into individual variables hard-codes number of items in select list.

```
rec  
company_pkg.allrows%ROWTYPE;  
BEGIN  
OPEN company_pkg.allrows;  
FETCH company_pkg.allrows INTO  
rec;  
IF rec.name = 'ACME' THEN ...  
CLOSE company_pkg.allrows;
```

If the cursor select list changes, it doesn't necessarily affect your code

- Complex data structures (collections, objects, records) can take up substantial amounts of memory
- Be aware of memory consumption
- Know how to analyze memory usage
- Adjust usage as needed
- Examine how you can do your own analysis

The BULK COLLECT Clause

- Use BULK COLLECT, % ATTRIBUTES wherever required
- BULK COLLECT loads multiple rows into the collections rather than one at a time
- Reduces strain on resources by reducing the pass to the database

```
CREATE OR REPLACE PROCEDURE
no_bulk_proc is
  CURSOR item_cur IS
    SELECT items.item_id, qty FROM items;

  item_rec item_cur%ROWTYPE;
BEGIN
  OPEN item_cur;
  LOOP
    FETCH item_cur INTO item_rec;
    EXIT WHEN item_cur%notfound;

    dbms_output.put_line(item_rec.item_id);
    dbms_output.put_line(item_rec.qty);
  END LOOP;
END no_bulk_proc;
```

```
CREATE OR REPLACE PROCEDURE bulk_proc IS
  CURSOR item_cur IS
    SELECT item_id, qty FROM items;

  TYPE t_item IS TABLE OF ITEMS.item_id%TYPE
  INDEX BY BINARY_INTEGER;
  TYPE t_qty IS TABLE OF ITEMS.qty%TYPE INDEX
  BY BINARY_INTEGER;

  v_item t_item;
  v_qty t_qty;
BEGIN
  OPEN item_cur;
  FETCH item_cur bulk collect INTO v_item, v_qty
  limit 100;

  FOR i IN v_item.first .. v_item.last LOOP
    dbms_output.put_line(v_item(i));
    dbms_output.put_line(v_qty(i));
  END LOOP;
  CLOSE item_cur;
END bulk_proc;
```


- Used for bulk-bind input operations before sending them to SQL engine
- Used with inserts, updates & deletes
- Used for Moving data from collections to tables
- Instead of executing repetitive, individual DML statements, you can write your code like this:

```
PROCEDURE upd_for_dept (...) IS
BEGIN
  FORALL indx IN list_of_emps.FIRST .. list_of_emps.LAST
    UPDATE employee
      SET salary = newsal_in
      WHERE employee_id = list_of_emps (indx);
END;
```

BEFORE Vs AFTER Trigger:


- Never use BEFORE trigger for Validations
- Use BEFORE triggers ONLY to modify: NEW value
- AFTER row triggers are slightly more efficient than BEFORE row triggers
- With BEFORE row triggers, affected data blocks must be read (logical read, not physical read) once for the trigger and then again for the triggering statement
- With AFTER row triggers, the data blocks must be read only once for both the triggering statement and the trigger

- SQL statements are among the most critical elements of our code base and should not be taken for granted
- SQL statements change constantly, forcing us to change code
- They cause most of the performance problems
- They cause many runtime errors in an application

- Don't repeat the same logical SQL statement
 - Repetition makes it almost impossible to maintain & optimize
- Use PL/SQL only when you need it
 - "Straight" SQL is almost always faster and easier
- Don't write any SQL statements in the application code (user-facing) layer
 - Both on the front end and the back end
 - That is, no hard-coded SQL in Java or .Net code



- Think of SQL as a service that is provided to you, not something you write
 - Or if you write it, you put it somewhere so that it can be easily found, reused and maintained
- This service has programs defined in the Data Access Layer
 - Known as table APIs or data Encapsulation, these programs have intelligence about business transactions & underlying tables



Tips to enhance performance in SQL scripts:

- Never do a calculation on an indexed column
(Example, WHERE salary*5 > :myvalue)
- Use the UNION statement instead of OR conditions
- Avoid the use of NOT IN or HAVING in the WHERE clause.
Instead, use the NOT EXISTS clause
- Specify numeric values in numeric form & character values
in character form

- Avoid specifying NULL in an indexed column
- Avoid the LIKE parameter if = will suffice. Using any Oracle function will invalidate the index, causing a full-table scan
- Never mix data types in Oracle queries, as it will invalidate the index. If the column is numeric, remember not to use quotes (e.g., salary = 50000). For char index columns, always use single quotes
- Oracle's rule-based optimizer looks at the order of table names in the FROM clause to determine the driving table
- Make sure that the last table specified in the FROM clause is the one that returns least rows
- In other words, specify multiple tables with the largest result set table specified first in the FROM clause

- Avoid using sub-queries when a JOIN will do the job
- Use the DECODE function to minimize the number of times a table has to be selected
- WHERE EXISTS sub-queries are better than JOIN if the number of records in driver query are less
- WHERE EXISTS can be better than JOIN when driving from parent records and we want to make sure that at least one child exists

- Try to group multiple sub queries into one
- Beware of Implicit data type conversions occurring on indexed columns
- NEVER use SELECT * FROM in application code
- Avoid using the DISTINCT clause as it always requires a sort operation
- Excessive use of DISTINCT clause may point to an underlying data model problem (e.g. a missing table)

- Application code must raise, handle, log and communicate errors in a consistent, robust manner
- Error Management - Sets the Rules
- Application developers do not write their own exception handling and raising code
 - Instead, they call programs to do the work for them
- Don't hard-code -20NNN error codes and their related messages
 - And if you use the -20NNN codes, maintain a repository of the numbers you have used
- Give names to any system exceptions you need to handle. Avoid codes like this:

```
WHEN OTHERS THEN  
    IF      SQLCODE = -24381 THEN...  
    ELSIF   SQLCODE = -1855  THEN...
```

- Use package-based utilities to raise, handle and log exceptions
- Compensate for PL/SQL weaknesses
- Some general guidelines:
 - Avoid hard-coding error numbers and messages
 - Build and use reusable components for raising, handling and logging errors
- Application-level code should not contain:
 - `RAISE_APPLICATION_ERROR`: Don't leave it to the developer to decide how to raise
 - `PRAGMA EXCEPTION_INIT`: Avoid duplication of error definitions

- CATG, BT16 team has written a basic document which describes the best coding practices to write SQL & PL/SQL



**Microsoft Office
Word 97 - 2003 Document**

In this session, we have covered:

- Best practices for PL/SQL Program
- Modularity
- Challenges of Local Module
- Learn Anchor Declarations of Variables
- Use of Bind Variables
- Cursors & Control Structures
- Bulk Collect Clause
- The FORALL Clause
- BEFORE Vs AFTER Trigger
- Optimizing SQL



Thank You